

Introduction

Vector databases store embeddings for fast similarity search.
They power retrieval for AI apps.

Benchmarks

Query latency dropped 42% after enabling HNSW indexes.

